

Date

# NoSQL

## Benefits of Relational Databases

- Following are the main benefits of Relational Databases
  - Designed for all purposes
  - ACID
  - Strong consistency, concurrency recovery
  - Mathematical background
  - Standard Query Language (SQL)
  - Lots of tools to use with i.e. Reporting services, entity frameworks,....

## ACID Properties

- To ensure integrity of data, the database system must maintain:

- Atomicity
  - \* Ensures that each transaction is treated as a single unit that either completes entirely or does not happen at all.
  - \* If one part of the transaction fails, the entire transaction fails and the database remains unchanged.

## • Consistency

- \* Ensures that a transaction takes the database from one valid state to another valid state, maintaining all defined rules, constraints, and data integrity.

## • Isolation

- \* Ensures that concurrent transactions do not interfere with each other.

- \* Each transaction should act as if it is the only one being executed, even if others are happening at the same time.

## • Durability

- \* Once a transaction is committed, the changes are permanent, even in the case of a system crash or power failure.

## Horizontal Scaling (Scaling Out)

- Horizontal scaling means that you scale by adding more machines into your pool of

Date

resources. Think of it like 2 more workers to get more work done.

## Vertical Scaling (Scaling Up)

→ Vertical scaling means that you scale by adding more power (CPU, RAM) to an existing machine.

## Impedance mismatch

→ An impedance mismatch can occur when accessing a relational database in an OOP language. Problems can arise because OOP languages like C++ or Python have very different approaches to accessing data.

→ Relational databases were not built for distributed applications because:

- Joins are expensive
- Hard to scale horizontally
- Impedance mismatch occurs
- Expensive (product cost, hardware, maintenance).

Also its weak in

- Speed (Performance)

- High availability
- Partition tolerance.

## NoSQL

→ A NoSQL database provides a mechanism for storage and retrieval of data that employs less constrained consistency models than traditional relational database.

→ NoSQL systems are also referred to as "Not only SQL" to emphasize that they do in fact allow SQL-like query languages to be used.

→ NoSQL avoids:

- Overhead of ACID transactions
- Complexity of SQL query
- Burden of up-front schema design
- DBA presence
- Transactions (it should be handled at application layer)

→ NoSQL provides:

- Easy and frequent changes to DB
- Fast development
- Large data volumes (Google)
- Schema Less

### Date

- When/Why to use NoSQL
  - Traditional RDBMS is too restrictive.
  - ACID support is not "really" needed.
  - Object-to-Relational (O/R) impedance
  - RDBMS is not distributed or scalable by nature
  - Logging data from distributed sources
  - Storing events / temporal data
  - Temporary Data
  - Data requiring flexible schema
  - Polyglot Persistence
    - \* Use different types of databases for different data needs.
    - \* NoSQL fits in when relational models don't suit a specific data type.
- When NOT to use NoSQL
  - Financial Data
  - Strict ACID compliance
  - Business-Critical Data

## Schema-Less Data-model

- In Relational Databases
  - You can't add a record which does not fit the schema.
  - You need to add NULLs to unused items in a row.
  - We should consider the datatypes i.e. you can't add a string to an integer field.
  - You can't add multiple items in a field (you should create another table: primary-key, foreign key, joins, normalization, ....).
- In NoSQL Databases:
  - There is no schema to consider.
  - There is no unused cell.
  - There is no datatype (implicit).
  - Most of considerations are done in application layer.
  - We gather all items in an aggregate (document).

## NoSQL Datamodels

- NoSQL databases are classified in four major datamodels:

## Date

- Key - value
  - Document
  - Column family
  - Graph
- Each database has its own query language.

### i) Key - value Data

- Key - value databases are most simplest NoSQL databases.
- Internally, these databases use a hash table structure where each unique key maps to a specific value.
- You use a key (string) to retrieve associated value.
- There is no schema or enforced structure for the stored values.
- All data are stored as pairs: one key , one value.
- Basic operations:
  - Insert (key, value)
  - Fetch (key)
  - Update (key)
  - Delete (key)

### ii) Column Family

- A column family is similar to a table in a relational database, but much more flexible:

- Data is stored in rows, and each row has a unique row key.
- Each row is made up of columns, but not all rows need to have the same columns.
- Columns are grouped into column families.
- Each column in a row is stored as a tuple: (column name, value, timestamp)

### Example:

→ Facebook search uses Cassandra which uses a column family data model. Here is a comparison b/w MySQL and Cassandra

- MySQL > 50 GB Data  
Writes Avg : ~ 300 ms  
Reads Avg : ~ 350 ms
- Cassandra > 50 GB Data  
Writes Avg : 0.12 ms  
Reads Avg : 15 ms

### iii) Graph Data Model

- The Graph Data Model is one of the core data models used in NoSQL databases. It is designed to represent and query highly

Date

connected data efficiently.

→ Following are main properties of Graph data model:

- Based on Graph Theory
- Scale vertically, no clustering
- You can use graph algorithms easily.
- Transactions
- ACID

#### iv) Document-Based data model

→ The Document-based data model is a popular type of NoSQL data model used by database like MongoDB.

→ It is designed to store, retrieve, and manage semi-structured data in the form of documents, usually in JSON, BSON or XML formats.

→ A document in this context is a self-contained unit of data similar to a record or row, but much more flexible and hierarchical.

→ Following are main properties of this data model:

- Each document is stored with a unique key. The

document can hold complex, structured data.

- Efficient data retrieval is achieved through B-Tree indexes. Indexes can be created on any fields within the documents.
- Documents can hold:
  - Key-value pairs
  - Key-array pairs
  - Nested documents

#### CAP Theorem

→ The CAP theorem states that it is impossible for any shared data-system to guarantee simultaneously all of the following three properties:

- Consistency — once data is written, all future read requests will contain that data.
  - Availability — the database is always available & responsive.
  - Partition Tolerance — if part of the database is unavailable, other parts are unaffected.
- We can't achieve all the three items in distributed database systems.



# Differences

	<b>SQL Databases</b>	<b>No SQL Database</b>
Example	Oracle , mysql	Mondo DB, CouchDB, Neo4J
Storage Model	Rows and tables	Key-value. Data stored as single document in JSON, XML
Schemas	Static	Dynamic
Scaling	Vertical & Horizontal	Horizontal
Transactions	Yes	Certain levels
Data Manipulation	Select, Insert , Update	Through Object Oriented API's

# In Conclusion!

- RDBMS is a great tool for solving ACID problems
  - When data validity is super important
  - When you need to support dynamic queries
- NoSQL is a great tool for solving data availability problems
  - When it's more important to have fast data than right data
  - When you need to scale based on changing requirements
- Pick the right tool for the job